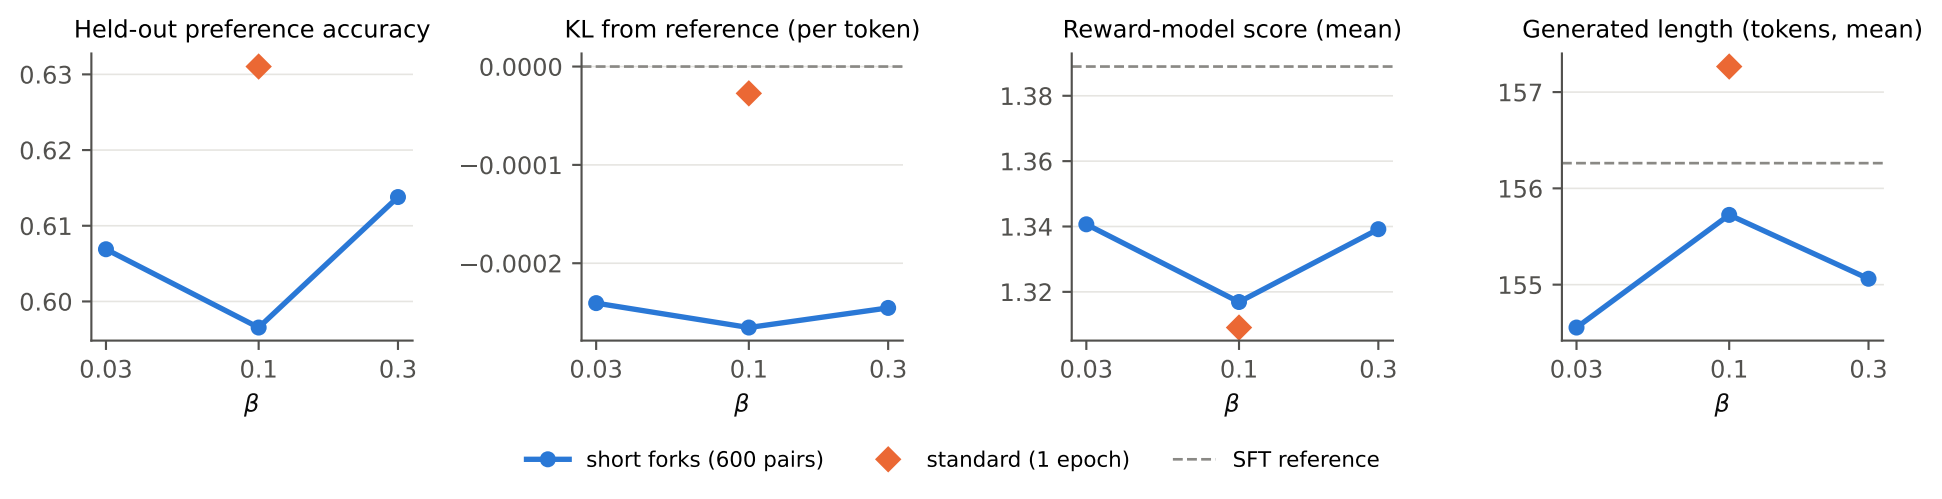
Held-out preference accuracy
0.63
0.62
0.61
0.60
0.03
0.1
0.3
β
KL from reference (per token)
0.0000
−0.0001
−0.0002
0.03
0.1
0.3
β
Reward-model score (mean)
1.38
1.36
1.34
1.32
0.03
0.1
0.3
β
Generated length (tokens, mean)
157
156
155
0.03
0.1
0.3
β
short forks (600 pairs)
standard (1 epoch)
SFT reference