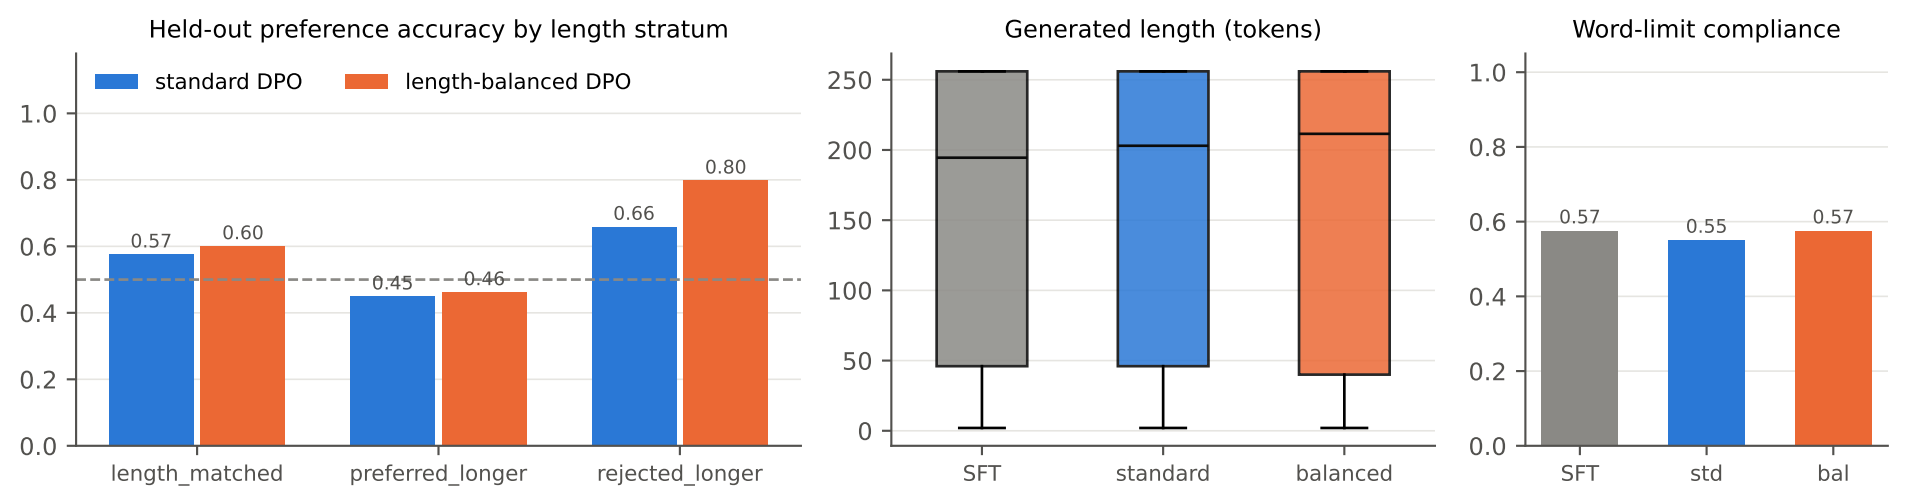
Held-out preference accuracy by length stratum
standard DPO
length-balanced DPO
1.0
0.8
0.6
0.4
0.2
0.0
0.57
0.60
0.45
0.46
0.66
0.80
length_matched
preferred_longer
rejected_longer
Generated length (tokens)
250
200
150
100
50
0
SFT
standard
balanced
Word-limit compliance
1.0
0.8
0.6
0.4
0.2
0.0
0.57
0.55
0.57
SFT
std
bal