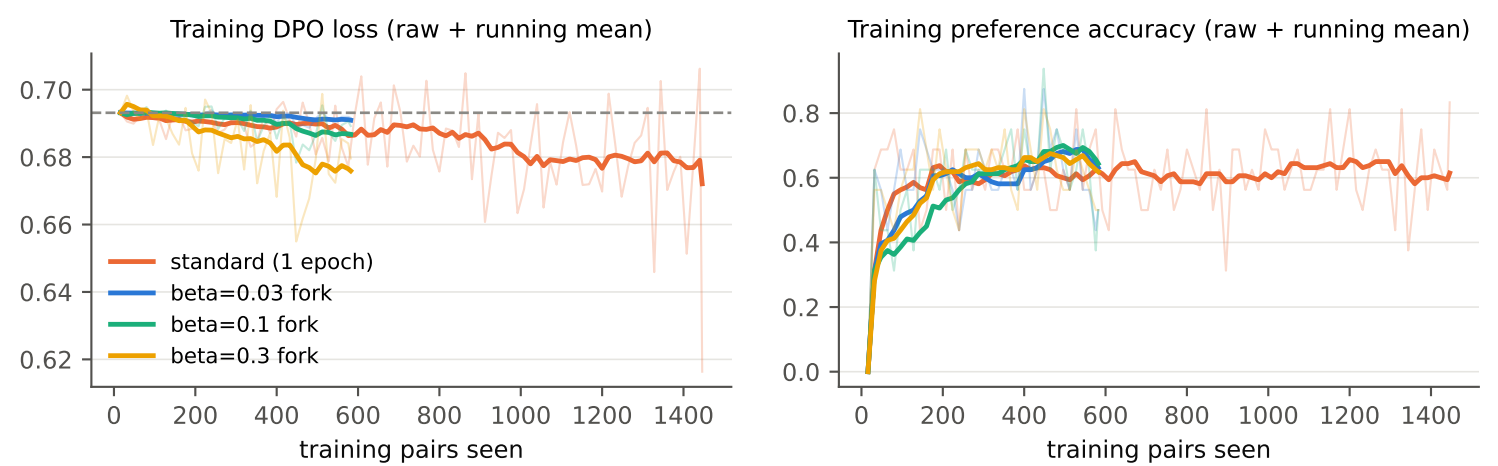
Training DPO loss (raw + running mean)
0.70
0.68
0.66
0.64
0.62
standard (1 epoch)
beta=0.03 fork
beta=0.1 fork
beta=0.3 fork
0
200
400
600
800
1000
1200
1400
training pairs seen
Training preference accuracy (raw + running mean)
0.8
0.6
0.4
0.2
0.0
0
200
400
600
800
1000
1200
1400
training pairs seen