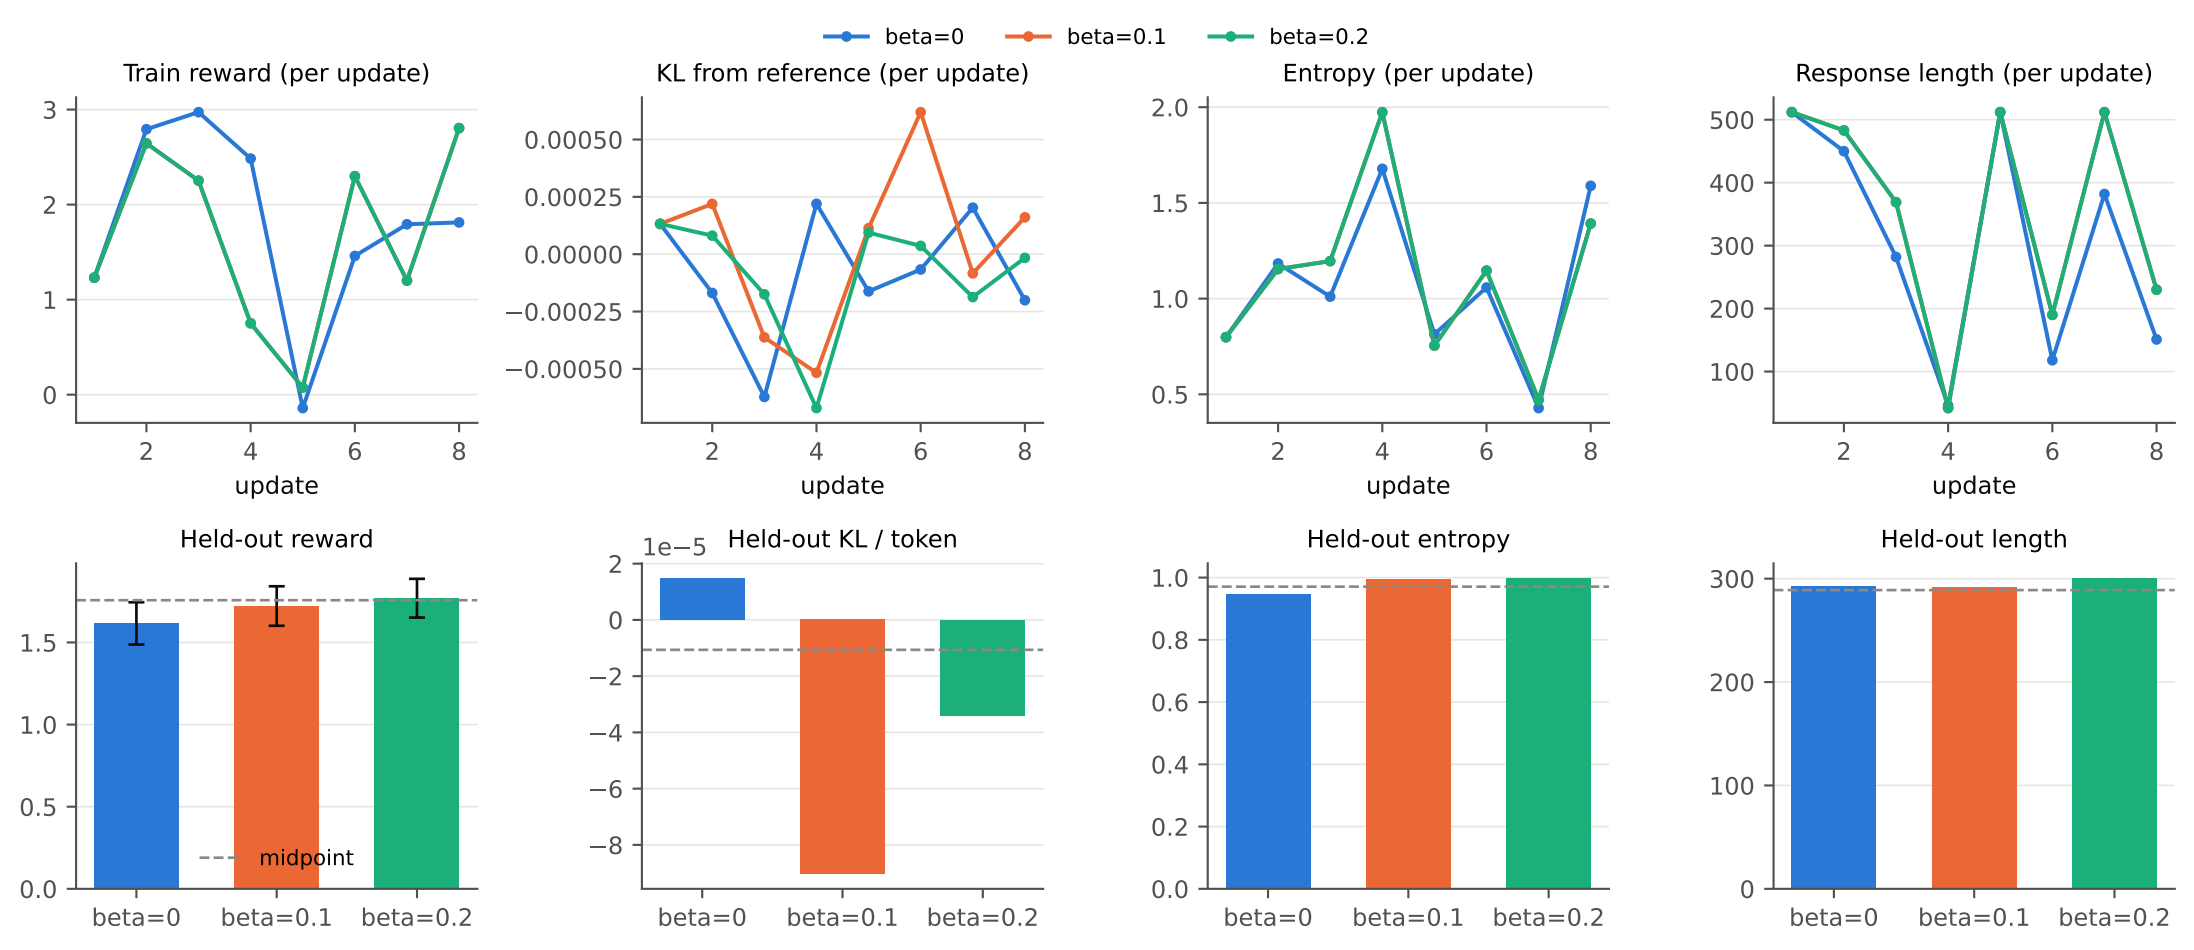
beta=0
beta=0.1
beta=0.2
Train reward (per update)
KL from reference (per update)
Entropy (per update)
Response length (per update)
update
Held-out reward
midpoint
Held-out KL / token
1e−5
Held-out entropy
Held-out length
beta=0
beta=0.1
beta=0.2