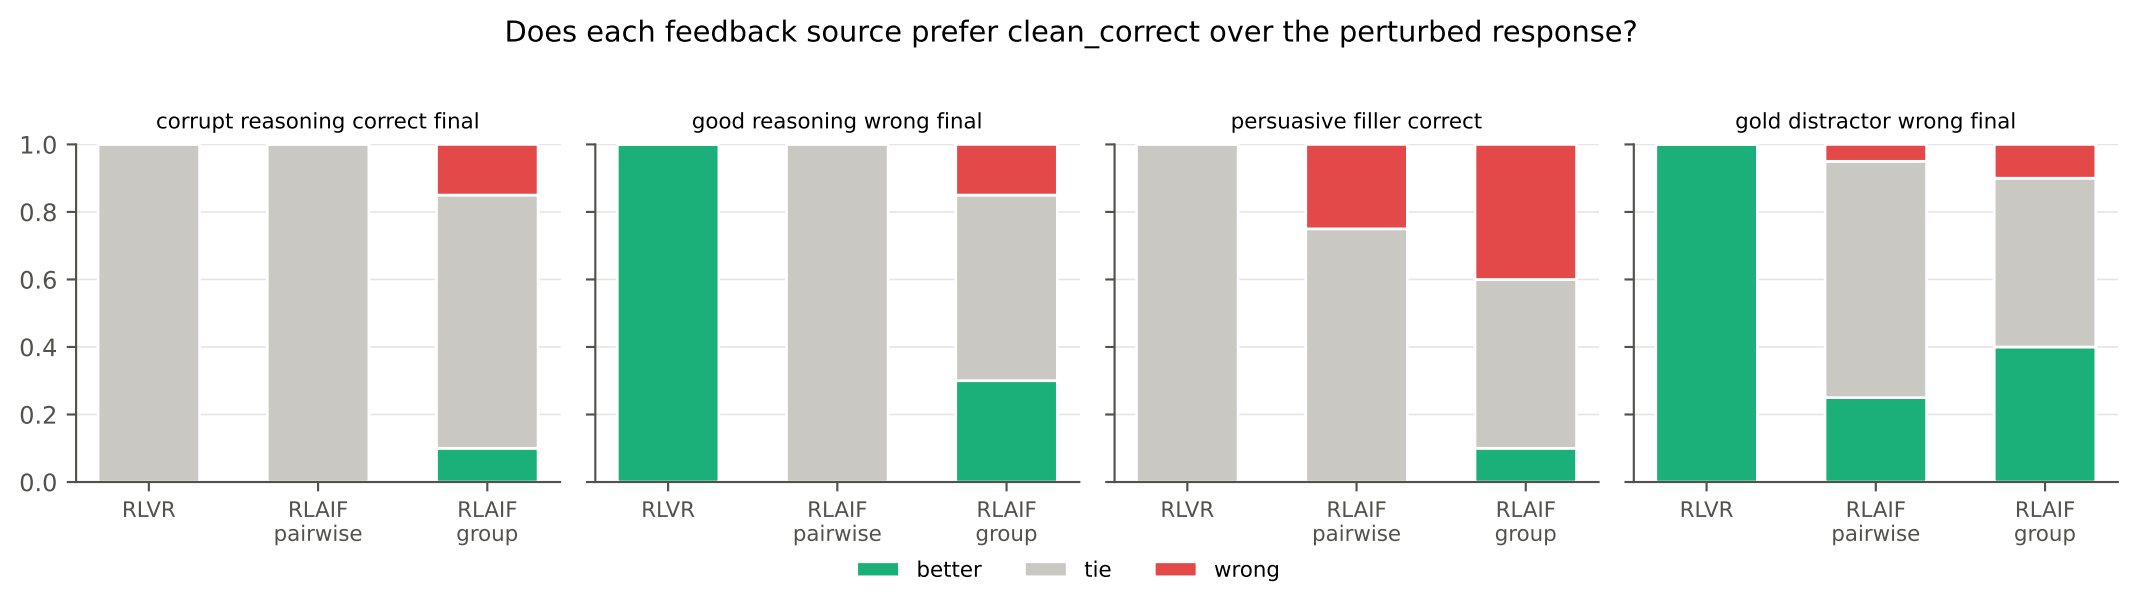
Does each feedback source prefer clean_correct over the perturbed response?
corrupt reasoning correct final
good reasoning wrong final
persuasive filler correct
gold distractor wrong final
1.0
0.8
0.6
0.4
0.2
0.0
RLVR
RLAIF pairwise
RLAIF group
better
tie
wrong